




Bench is selected for DB disagreement and is NOT a random topological sample: GPI-anchored / lipidated / single-pass classes are over-represented (DBs disagree on these); multi-pass TM is under-represented (DBs agree on these). Bench-accuracy is therefore a lower bound on full-universe accuracy. Error bars: Wilson 95% binomial CI.